

ARJUN MEHTA

arjun.mehta@gmail.com | +91 98765 43210 | linkedin.com/in/arjunmehta | github.com/arjunmehta-ml

PROFESSIONAL SUMMARY

Senior Machine Learning Engineer with 5 years of experience building and deploying production ML systems at scale. Specializes in end-to-end ML pipelines from data ingestion through model deployment using FastAPI and batch processing frameworks. Proven track record of delivering high-impact ML solutions across fintech and e-commerce domains.

TECHNICAL SKILLS

Languages: Python (expert), R, Scala

ML/DL Frameworks: scikit-learn, TensorFlow, PyTorch, XGBoost, LightGBM

Model Deployment: FastAPI, Flask, Docker, Kubernetes, AWS SageMaker, batch inference pipelines

Big Data: Apache Spark, Hadoop, Hive, processing datasets of 10B+ records

MLOps: MLflow, Airflow, Kubeflow, CI/CD for ML pipelines

Cloud: AWS (EC2, S3, Lambda, SageMaker), GCP (BigQuery, Vertex AI)

PROFESSIONAL EXPERIENCE

Senior ML Engineer — Razorpay, Bengaluru | Jan 2022 – Present

- * Built fraud detection model using gradient boosting (XGBoost) achieving 94% precision on 500M+ daily transactions — deployed as FastAPI microservice handling 50K RPS
- * Designed batch inference pipeline using Apache Airflow processing 2TB of transaction data nightly for risk scoring
- * Implemented automated retraining pipeline using MLflow tracking, reducing model drift incidents by 60%
- * Led migration of legacy scoring models to containerized FastAPI services, reducing latency from 450ms to 28ms
- * Trained and fine-tuned NLP models (BERT-based) for merchant category classification across 15M merchant profiles

ML Engineer — Flipkart, Bengaluru | Jul 2019 – Dec 2021

- * Built product recommendation engine using collaborative filtering and matrix factorization on 100M+ user-item interactions, deployed via Flask REST API serving 8M daily active users
- * Developed demand forecasting models (LSTM, ARIMA, Prophet) for 5M SKUs across 200+ categories using Spark on 3 years of historical sales data
- * Implemented A/B testing framework to evaluate model variants, managed statistical significance testing across concurrent experiments
- * Built real-time feature store using Redis and Kafka for online model serving with sub-10ms feature retrieval

KEY PROJECTS

Credit Risk Scoring System — End-to-end ML pipeline for personal loan default prediction. Trained logistic regression, random forest, and gradient boosting ensemble on 50M loan records. Deployed as Flask microservice integrated with credit bureau APIs. AUC-ROC: 0.91.

Real-Time Anomaly Detection — Built LSTM autoencoder for detecting anomalous patterns in IoT sensor streams. Processes 100K events/sec via Kafka. Model served via FastAPI with async batch prediction endpoints.

NLP Document Classifier — Fine-tuned DistilBERT for multi-label document classification across 50 categories. Dataset: 10M documents. Deployed via FastAPI with model versioning and shadow deployment support.

EDUCATION

B.Tech, Computer Science — IIT Bombay, 2019 | CGPA: 8.7/10

Relevant coursework: Machine Learning, Deep Learning, Statistical Inference, Algorithms, Distributed Systems

CERTIFICATIONS

AWS Certified Machine Learning Specialty (2023) | Google Professional ML Engineer (2022) | Deep Learning Specialization — Coursera/deeplearning.ai